

```
#Import necessary Libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
%matplotlib inline

plt.style.use('dark_background')

import seaborn as sns
from nltk import word_tokenize

from sklearn.model_selection import train_test_split
from sklearn.preprocessing import LabelEncoder
from sklearn.metrics import auc, roc_auc_score, roc_curve, confusion_matrix, classification_r

from keras.models import Model, load_model
from keras.layers import LSTM, Activation, Dense, Dropout, Input, Embedding
from keras.optimizers import RMSprop
from keras.preprocessing.text import Tokenizer
from tensorflow.keras.preprocessing.sequence import pad_sequences
from keras.utils import to_categorical
from keras.callbacks import plot_model, EarlyStopping, ModelCheckpoint, TensorBoard
nltk.download('stopwords')
nltk.download('wordnet')
nltk.download('punkt')
nltk.download('omw-1.4')
```



```
[nltk_data] Downloading package stopwords to /root/nltk_data...
[nltk_data] Package stopwords is already up-to-date!
[nltk_data] Downloading package wordnet to /root/nltk_data...
[nltk_data] Package wordnet is already up-to-date!
[nltk_data] Downloading package punkt to /root/nltk_data...
[nltk_data] Package punkt is already up-to-date!
[nltk_data] Downloading package omw-1.4 to /root/nltk_data...
[nltk_data] Package omw-1.4 is already up-to-date! True
```

```
df=pd.read_csv("/content/drive/MyDrive/IBM/Assignment - 4/spam.csv",encoding="latin-1");
df.head()
```

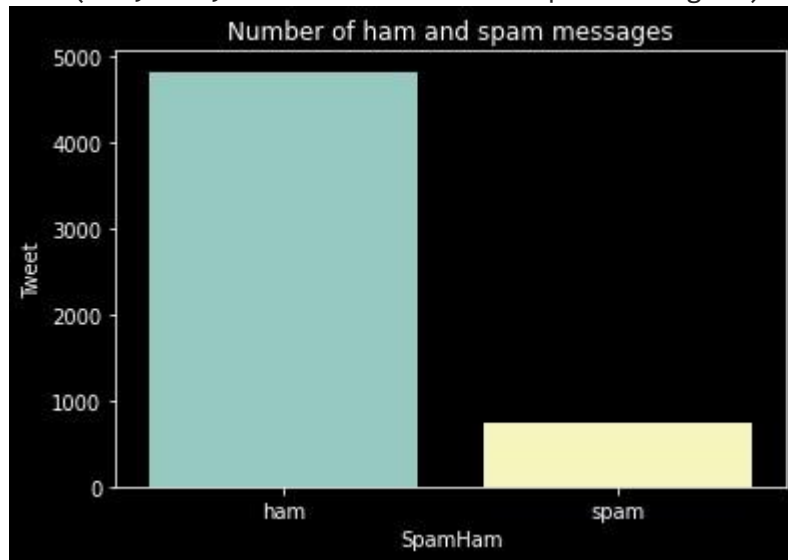
	v1	v2	Unnamed: 2	Unnamed: 3	Unnamed: 4
0	ham	Go until jurong point, crazy.. Available only ...	NaN	NaN	NaN
1	ham	Ok lar... Joking wif u oni...	NaN	NaN	NaN
2	spam	Free entry in 2 a wkly comp to win FA Cup fina...	NaN	NaN	NaN
3	ham	U dun say so early hor... U c already then say...	NaN	NaN	NaN

```
df.shape
```

```
(5572, 5)
```

```
fig, ax = plt.subplots()
sns.countplot(df.v1, ax=ax)
ax.set_xlabel('SpamHam')
ax.set_ylabel('Tweet')
ax.set_title('Number of ham and spam messages')
```

Text(0.5, 1.0, 'Number of ham and spam messages')



```
X = df.loc[:, 'v2']
y = df.loc[:, 'v1']
```

```
X
0          Go until jurong point, crazy.. Available only ...
1          Ok lar... Joking wif u oni...
2          Free entry in 2 a wkly comp to win FA Cup fina...
3          U dun say so early hor... U c already then say...
4          Nah I don't think he goes to usf, he lives aro...
...
5567       This is the 2nd time we have tried 2 contact u...
5568       Will i_b going to esplanade fr home?
5569       Pity, * was in mood for that. So...any other s...
5570       The guy did some bitching but I acted like i'd...
5571       Rofl. Its true to its name Name: v2, Length: 5572, dtype: object
```

```
X_train_data, X_test_data, y_train_labels, y_test_labels = train_test_split(X, y, test_size=0)
print(X_train_data.shape) print(X_test_data.shape)
```


```
(4457,)
(1115,)
sent_lens = [] for sent
in X_train_data:
    sent_lens.append(len(word_tokenize(sent)))

print(max(sent_lens))
```

220

```
sns.distplot(sent_lens, bins=10, kde=True)
```

```
<matplotlib.axes._subplots.AxesSubplot at 0x7f81b671abd0>
```



```
np.quantile(sent_lens, 0.95)
```

```
39.0
```

```
max_sequence_length = 38
```

```
tok = Tokenizer()
tok.fit_on_texts(X_train_data.values)
```

```
vocab_length = len(tok.word_index) #len(tok.word_counts) or len(tok.index_word.keys()) will a
print('No. of unique tokens(vocab_size): ', vocab_length)
```

```
X_train_sequences = tok.texts_to_sequences(X_train_data.values) X_test_sequences =
tok.texts_to_sequences(X_test_data.values) print('No of sequences:', len(X_train_sequences))
#No of sequences will be same as the number print(X_train_sequences[:2])
```

```
#make all sequences of equal length
X_train = pad_sequences(X_train_sequences, maxlen=max_sequence_length)
X_test = pad_sequences(X_test_sequences, maxlen=max_sequence_length)
X_train[:2]
```

```
No. of unique tokens(vocab_size): 7954
```

```
No of sequences: 4457
```

```
[[38, 30, 8, 5, 273, 1989, 81, 116, 26, 11, 1656, 322, 10, 53, 18, 299, 30, 349, 1990],
array([[ 0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0],
       [ 0,  0,  0,  0,  0,  0,  0, 38, 30,  8,  5, 273, 1989, 81, 116, 26, 11, 1656, 322, 10, 53,
        18, 299, 30, 349, 1990],
       [ 0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,
        0,  0,  0,  0, 799, 15, 2555, 1442, 1127, 192, 2556,
```

```
171, 12, 98, 1991, 44, 195, 1657, 2557, 1992, 2558, 21,
9, 4, 203, 1025, 225]], dtype=int32)
```

```
y_train_labels.values array(['ham', 'spam', 'ham', ..., 'ham', 'ham',
'ham'], dtype=object)
```

```
le = LabelEncoder() y_train =
le.fit_transform(y_train_labels) y_test =
le.fit_transform(y_test_labels)
print(y_train)
```

```
[0 1 0 ... 0 0 0]
```

```
def create_model(vocab_len, max_seq_len):
    inputs = Input(name='inputs', shape=[max_seq_len]) #None, 150    layer =
    Embedding(vocab_length + 1, 50, input_length=max_seq_len)(inputs) #None, 150, 50    layer
    = LSTM(64)(layer) #None, 64    layer = Dense(256,name='FC1')(layer) #None, 256    layer
    = Activation('relu')(layer) #None, 256    layer = Dropout(0.5)(layer) #None, 256    layer
    = Dense(1,name='out_layer')(layer) #None, 1    layer = Activation('sigmoid')(layer) #None,
    1    model = Model(inputs=inputs,outputs=layer)
    model.compile(loss='binary_crossentropy',optimizer=RMSprop(), metrics=['acc'])    return
    model
```

```
model = create_model(vocab_length, max_sequence_length)
model.summary()
```

```
Model: "model_1"
```

Layer (type)	Output Shape	Param #
=====		
inputs (InputLayer)	[(None, 38)]	0
embedding_1 (Embedding)	(None, 38, 50)	397750
lstm_1 (LSTM)	(None, 64)	29440
FC1 (Dense)	(None, 256)	16640
activation_2 (Activation)	(None, 256)	0
dropout_1 (Dropout)	(None, 256)	0
out_layer (Dense)	(None, 1)	257
activation_3 (Activation)	(None, 1)	0
=====		
Total params: 444,087		
Trainable params: 444,087		
Non-trainable params: 0		

```
filepath='model_with_best_weights.h5'
callbacks = [EarlyStopping(monitor='val_loss', patience=5, verbose=1),
             ModelCheckpoint(filepath=filepath, monitor='val_loss', save_best_only=True, verb
```

]

```
history = model.fit(X_train, y_train, batch_size=128, epochs=20, validation_split=0.2, callba
```

```
Epoch 1/20
28/28 [=====] - ETA: 0s - loss: 0.2985 - acc: 0.8853
Epoch 1: val_loss improved from inf to 0.13599, saving model to model_with_best_weights
28/28 [=====] - 6s 99ms/step - loss: 0.2985 - acc: 0.8853 - val
Epoch 2/20
28/28 [=====] - ETA: 0s - loss: 0.0737 - acc: 0.9832
Epoch 2: val_loss improved from 0.13599 to 0.05181, saving model to model_with_best_weig
28/28 [=====] - 2s 76ms/step - loss: 0.0737 - acc: 0.9832 - val
Epoch 3/20
28/28 [=====] - ETA: 0s - loss: 0.0318 - acc: 0.9924
Epoch 3: val_loss did not improve from 0.05181
28/28 [=====] - 2s 75ms/step - loss: 0.0318 - acc: 0.9924 - val
Epoch 4/20
28/28 [=====] - ETA: 0s - loss: 0.0160 - acc: 0.9950
Epoch 4: val_loss improved from 0.05181 to 0.04797, saving model to model_with_best_weig
28/28 [=====] - 2s 78ms/step - loss: 0.0160 - acc: 0.9950 - val
Epoch 5/20
28/28 [=====] - ETA: 0s - loss: 0.0083 - acc: 0.9975
Epoch 5: val_loss did not improve from 0.04797
28/28 [=====] - 2s 74ms/step - loss: 0.0083 - acc: 0.9975 - val
Epoch 6/20
28/28 [=====] - ETA: 0s - loss: 0.0039 - acc: 0.9986
Epoch 6: val_loss did not improve from 0.04797
28/28 [=====] - 2s 74ms/step - loss: 0.0039 - acc: 0.9986 - val
Epoch 7/20
28/28 [=====] - ETA: 0s - loss: 0.0014 - acc: 0.9997
Epoch 7: val_loss did not improve from 0.04797
28/28 [=====] - 2s 74ms/step - loss: 0.0014 - acc: 0.9997 - val
Epoch 8/20
28/28 [=====] - ETA: 0s - loss: 0.0012 - acc: 0.9997
Epoch 8: val_loss did not improve from 0.04797
28/28 [=====] - 2s 75ms/step - loss: 0.0012 - acc: 0.9997 - val
Epoch 9/20
28/28 [=====] - ETA: 0s - loss: 4.6690e-04 - acc: 0.9997
Epoch 9: val_loss did not improve from 0.04797
28/28 [=====] - 2s 75ms/step - loss: 4.6690e-04 - acc: 0.9997
Epoch 9: early stopping
```


```
history_dict = history.history
```

```
# list all data in history
print(history_dict.keys())
```


```
# summarize history for loss
plt.plot(history_dict['loss'])
plt.plot(history_dict['val_loss'])
plt.title('Training and Validation Loss')
```

```
plt.ylabel('loss') plt.xlabel('epoch')
plt.legend(['train', 'test'], loc='upper
left') plt.show()
```

```
# summarize history for accuracy
plt.plot(history_dict['acc'])
plt.plot(history_dict['val_acc'])
plt.title('Training and Validation Accuracy')
plt.ylabel('accuracy') plt.xlabel('epoch')
plt.legend(['train', 'test'], loc='upper
left') plt.show()
dict_keys(['loss', 'acc', 'val_loss', 'val_acc'])
y_pred
```



```
loaded_model = load_model('model_with_best_weights.h5')
test_loss, test_acc = accr = loaded_model.evaluate(X_test, y_test)
print('test set\n Loss: {:.03f}\n Accuracy: {:.03f}'.format(test_loss, test_acc))
35/35 [====] - 15.9ms/step - loss: 0.0634 - acc: 0.9839
Test set
Loss: 0.063
Accuracy: 0.983
```



```
y_pred_proba = loaded_model.predict(X_test)
# y_pred_proba = loaded_model.predict_classes(X_test) #we can't use it on Model object. Can be use
print(y_pred_proba.round(y_pred_proba, 3))
y_pred = y_pred_proba > 0.5
```

```
35/35 [====] - 2s 8ms/step
[[0.644]
 [0.837]
 ...
 [0.
 [0.001]
 [0.987]]
array([[ True],
 [False],
 [ True],
 ...,
 [False],
 [False],
 [ True]])
```

```
# summarize the first few cases for i in range(5):    print('%s => %d
(expected %d)' % (X_test[i].tolist(), y_pred[i], y_test[i]))
```

```
[0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1366, 1578, 1432, 19, 78
[1, 188, 11, 6440, 2, 7, 1, 135, 2, 28, 12, 4, 290, 7931, 1, 104, 33, 3, 22, 647, 15, 28
[0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 39, 54, 258, 144, 3, 54, 21
[0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 64, 33, 3, 1528, 13, 263, 53, 79, 228, 79, 3, 31, 7, 838,
[0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 715, 29, 357, 532, 622, 15, 1107, 528, 706, 4
```

```
print(confusion_matrix(y_test, y_pred))
```

```
[[964   1]
 [ 17 133]]
```

```
print(classification_report(y_test, y_pred))
```

```
precision    recall  f1-score   support

         0           0.98         1.00         0.99         965
         1           0.99         0.89         0.94        150

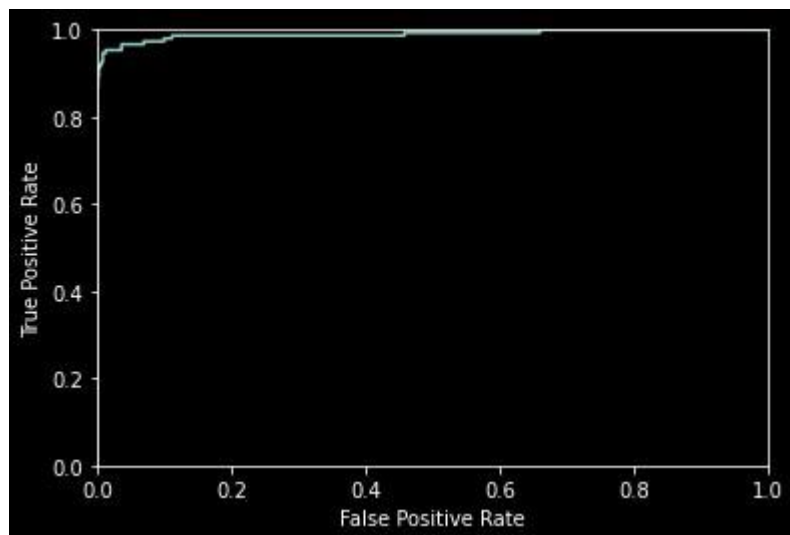
 accuracy
macro avg           0.99         0.94         0.96        1115 weighted
avg           0.98         0.98         0.98        1115
```

```
auc = roc_auc_score(y_test, y_pred_proba)
print('AUC: %.3f' % auc)
```

```
AUC: 0.990
```

```
fpr_keras, tpr_keras, thresholds_keras = roc_curve(y_test, y_pred_proba)
```

```
def plot_roc_curve(fpr, tpr):
import matplotlib.pyplot as plt
plt.plot(fpr, tpr)
plt.axis([0,1,0,1])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.show()
plot_roc_curve (fpr_keras,
tpr_keras)
```



[Colab paid products - Cancel contracts here](#)

